cowrie Documentation

Release 20.3.0

Michel Oosterhof

Contents:

Cowrie

1.1 Welcome to the Cowrie GitHub repository

This is the official repository for the Cowrie SSH and Telnet Honeypot effort.

1.2 What is Cowrie

Cowrie is a medium to high interaction SSH and Telnet honeypot designed to log brute force attacks and the shell interaction performed by the attacker. In medium interaction mode (shell) it emulates a UNIX system in Python, in high interaction mode (proxy) it functions as an SSH and telnet proxy to observe attacker behavior to another system.

Cowrie is maintained by Michel Oosterhof.

1.3 Documentation

The Documentation can be found here.

1.4 Slack

You can join the Cowrie community at the following Slack workspace.

1.5 Features

- **Choose to run as an emulated shell (default):**

  – Fake filesystem with the ability to add/remove files. A full fake filesystem resembling a Debian 5.0 installation is included

  – Possibility of adding fake file contents so the attacker can *cat* files such as */etc/passwd*. Only minimal file contents are included

  – Cowrie saves files downloaded with wget/curl or uploaded with SFTP and scp for later inspection

- **Or proxy SSH and telnet to another system**

  – Run as a pure telnet and ssh proxy with monitoring

  – Or let Cowrie manage a pool of Qemu emualted servers to provide the systems to login to

For both settings:

- Session logs are stored in an UML Compatible format for easy replay with the *bin/playlog* utility.

- SFTP and SCP support for file upload

- Support for SSH exec commands

- Logging of direct-tcp connection attempts (ssh proxying)

- Forward SMTP connections to SMTP Honeypot (e.g. mailoney)

- JSON logging for easy processing in log management solutions

1.6 Docker

Docker versions are available.

- To get started quickly and give Cowrie a try, run:

```
docker run -p 2222:2222 cowrie/cowrie
ssh -p 2222 root@localhost
```

- On Docker Hub: https://hub.docker.com/r/cowrie/cowrie

- Or get the Dockerfile directly at https://github.com/cowrie/docker-cowrie

1.7 Requirements

Software required:

- Python 3.5+

- python-virtualenv

For Python dependencies, see requirements.txt.

1.8 Files of interest:

- *etc/cowrie.cfg* - Cowrie's configuration file. Default values can be found in etc/cowrie.cfg.dist.
- *share/cowrie/fs.pickle* - fake filesystem
- *etc/userdb.txt* - credentials to access the honeypot
- honeyfs/ - file contents for the fake filesystem - feel free to copy a real system here or use *bin/fsctl*
- *honeyfs/etc/issue.net* - pre-login banner
- honeyfs/etc/motd - post-login banner
- *var/log/cowrie/cowrie.json* - transaction output in JSON format
- *var/log/cowrie/cowrie.log* - log/debug output
- *var/lib/cowrie/tty/* - session logs, replayable with the *bin/playlog* utility.
- *var/lib/cowrie/downloads/* - files transferred from the attacker to the honeypot are stored here
- share/cowrie/txtcmds/ - file contents for simple fake commands
- bin/createfs - used to create the fake filesystem
- bin/playlog - utility to replay session logs

1.9 Data Sharing

Cowrie will by default upload data on crashes and Python exceptions to api.cowrie.org. This information is used to improve the honeypot and is not shared with third parties. It can be disabled by setting *enabled=false* in *[output_crashreporter]*.

1.10 Contributors

Many people have contributed to Cowrie over the years. Special thanks to:

- Upi Tamminen (desaster) for all his work developing Kippo on which Cowrie was based
- Dave Germiquet (davegermiquet) for TFTP support, unit tests, new process handling
- Olivier Bilodeau (obilodeau) for Telnet support
- Ivan Korolev (fe7ch) for many improvements over the years.
- Florian Pelgrim (craneworks) for his work on code cleanup and Docker.
- Guilherme Borges (sgtpepperpt) for SSH and telnet proxy (GSoC 2019)
- And many many others.

Frequently asked questions

2.1 Why can't I start cowrie on port 22?

The possible answer for that is you might already have a service(possibly SSH) running on that port so setting up Cowrie on that port will cause a problem. Try changing the port in *listen_endpoints* present in config file(cowrie.cfg.dist/cowrie.cfg).

2.2 Why do I get logged into my own system when accessing cowrie on port 22?

This is probably a similar problem as it was in the above question. This can also be fixed by changing the port in the config file.

2.3 What I am getting permission denied when running cowrie on port 22?

You need root privileges to run Cowrie on any port lower than 1024. This can be fixed by setting up Authbind.

2.4 Do I need to copy all the content of cowrie.cfg.dist to cowrie.cfg?

No, Cowrie can read add your only local changes to cowrie.cfg and the remaining settings will automatically be read from cowrie.cfg.dist

2.5 Why certain commands aren't implemented?

There are lots of UNIX command implemented in cowrie and that is because Cowrie is more focused to provide proxy support i.e use Cowrie to connect to an actual machine that is actual machine having support for all the UNIX functionalities.

LICENSE

Installing Cowrie in seven steps.

This guide describes how to install Cowrie in *shell* mode. For *proxy* mode read *PROXY.rst*.

- [Step 1: Install dependencies](#step-1-install-dependencies)

- [Step 2: Create a user account](#step-2-create-a-user-account)

- [Step 3: Checkout the code](#step-3-checkout-the-code)

- [Step 4: Setup Virtual Environment](#step-4-setup-virtual-environment)

- [Step 5: Install configuration file](#step-5-install-configuration-file)

- [Step 6: Generate a DSA key (OPTIONAL)](#step-6-generate-a-dsa-key)

- [Step 7: Starting Cowrie](#step-7-turning-on-cowrie)

- [Step 8: Port redirection (OPTIONAL)](#step-8-port-redirection-optional)

- [Running within supervisord (OPTIONAL)](#running-using-supervisord)

- [Configure Additional Output Plugins (OPTIONAL)](#configure-additional-output-plugins-optional)

- [Troubleshooting](#troubleshooting)

4.1 Step 1: Install dependencies

First we install system-wide support for Python virtual environments and other dependencies. Actual Python packages are installed later.

On Debian based systems (last verified on Debian 10, 2019-08-18): For a Python3 based environment:

```
$ sudo apt-get install git python-virtualenv libssl-dev libffi-dev build-essential␣
→libpython3-dev python3-minimal authbind virtualenv
```

4.2 Step 2: Create a user account

It's strongly recommended to run with a dedicated non-root user id:

```
$ sudo adduser --disabled-password cowrie
Adding user 'cowrie' ...
Adding new group 'cowrie' (1002) ...
Adding new user 'cowrie' (1002) with group 'cowrie' ...
Changing the user information for cowrie
Enter the new value, or press ENTER for the default
Full Name []:
Room Number []:
Work Phone []:
Home Phone []:
Other []:
Is the information correct? [Y/n]

$ sudo su - cowrie
```

4.3 Step 3: Checkout the code

Check out the code:

```
$ git clone http://github.com/cowrie/cowrie
Cloning into 'cowrie'...
remote: Counting objects: 2965, done.
remote: Compressing objects: 100% (1025/1025), done.
remote: Total 2965 (delta 1908), reused 2962 (delta 1905), pack-reused 0
Receiving objects: 100% (2965/2965), 3.41 MiB | 2.57 MiB/s, done.
Resolving deltas: 100% (1908/1908), done.
Checking connectivity... done.

$ cd cowrie
```

4.4 Step 4: Setup Virtual Environment

Next you need to create your virtual environment:

```
$ pwd
/home/cowrie/cowrie
$ virtualenv --python=python3 cowrie-env
New python executable in ./cowrie/cowrie-env/bin/python
Installing setuptools, pip, wheel...done.
```

Activate the virtual environment and install packages:

```
$ source cowrie-env/bin/activate
(cowrie-env) $ pip install --upgrade pip
(cowrie-env) $ pip install --upgrade -r requirements.txt
```

4.5 Step 5: Install configuration file

The configuration for Cowrie is stored in cowrie.cfg.dist and cowrie.cfg. Both files are read on startup, where entries from cowrie.cfg take precedence. The .dist file can be overwritten by upgrades, cowrie.cfg will not be touched. To run with a standard configuration, there is no need to change anything. To enable telnet, for example, create cowrie.cfg and input only the following:

```
[telnet]
enabled = true
```

4.6 Step 6: Starting Cowrie

Start Cowrie with the cowrie command. You can add the cowrie/bin directory to your path if desired. An existing virtual environment is preserved if activated, otherwise Cowrie will attempt to load the environment called "cowrie-env":

```
$ bin/cowrie start
Activating virtualenv "cowrie-env"
Starting cowrie with extra arguments [] ...
```

4.7 Step 7: Listening on port 22 (OPTIONAL)

There are three methods to make Cowrie accessible on the default SSH port (22): *iptables*, *authbind* and *setcap*.

4.7.1 Iptables

Port redirection commands are system-wide and need to be executed as root. A firewall redirect can make your existing SSH server unreachable, remember to move the existing server to a different port number first.

The following firewall rule will forward incoming traffic on port 22 to port 2222 on Linux:

```
$ sudo iptables -t nat -A PREROUTING -p tcp --dport 22 -j REDIRECT --to-port 2222
```

Or for telnet:

```
$ sudo iptables -t nat -A PREROUTING -p tcp --dport 23 -j REDIRECT --to-port 2223
```

Note that you should test this rule only from another host; it doesn't apply to loopback connections.

On MacOS run:

```
$ echo "rdr pass inet proto tcp from any to any port 22 -> 127.0.0.1 port 2222" |␣
→sudo pfctl -ef -
```

4.7.2 Authbind

Alternatively you can run authbind to listen as non-root on port 22 directly:

```
$ sudo apt-get install authbind
$ sudo touch /etc/authbind/byport/22
$ sudo chown cowrie:cowrie /etc/authbind/byport/22
$ sudo chmod 770 /etc/authbind/byport/22
```

Edit bin/cowrie and modify the AUTHBIND_ENABLED setting

Change the listening port to 22 in cowrie.cfg:

```
[ssh]
listen_endpoints = tcp:22:interface=0.0.0.0
```

Or for telnet:

```
$ apt-get install authbind
$ sudo touch /etc/authbind/byport/23
$ sudo chown cowrie:cowrie /etc/authbind/byport/23
$ sudo chmod 770 /etc/authbind/byport/23
```

Change the listening port to 23 in cowrie.cfg:

```
[telnet]
listen_endpoints = tcp:2223:interface=0.0.0.0
```

4.7.3 Setcap

Or use setcap to give permissions to Python to listen on ports<1024:

```
$ setcap cap_net_bind_service=+ep /usr/bin/python3
```

And change the listening ports in *cowrie.cfg* as above.

4.8 Installing Backend Pool dependencies (OPTIONAL)

If you want to use the proxy functionality combined with the automatic backend pool, you need to install some dependencies, namely qemu, libvirt, and their Python interface. In Debian/Ubuntu:

```
$ sudo apt-get install qemu qemu-system-arm qemu-system-x86 libvirt-dev libvirt-
→daemon libvirt-daemon-system libvirt-clients nmap
```

Then install the Python API to run the backend pool:

```
(cowrie-env) $ pip install libvirt-python==5.5.0
```

To allow Qemu to use disk images and snapshots, set it to run with the user and group of the user running the pool (usually called 'cowrie' too:

```
$ sudo vim /etc/libvirt/qemu.conf
```

Search and set both *user* and *group* to *"cowrie"*, or the username/group you'll be running the backend pool with.

4.9 Running using Supervisord (OPTIONAL)

On Debian, put the below in /etc/supervisor/conf.d/cowrie.conf:

```
[program:cowrie]
command=/home/cowrie/cowrie/bin/cowrie start
directory=/home/cowrie/cowrie/
user=cowrie
autorestart=true
redirect_stderr=true
```

Update the bin/cowrie script, change:

```
DAEMONIZE=""
```

to:

```
DAEMONIZE="-n"
```

4.10 Configure Additional Output Plugins (OPTIONAL)

Cowrie automatically outputs event data to text and JSON log files in *var/log/cowrie*. Additional output plugins can be configured to record the data other ways. Supported output plugins include:

- Cuckoo
- ELK (Elastic) Stack
- Graylog
- Kippo-Graph
- Splunk
- SQL (MySQL, SQLite3, RethinkDB)

See ~/cowrie/docs/[Output Plugin]/README.rst for details.

Troubleshooting

If you see *twistd: Unknown command: cowrie* **there are two** possibilities. If there's a Python stack trace, it probably means there's a missing or broken dependency. If there's no stack trace, double check that your PYTHON-PATH is set to the source code directory.

Default file permissions

To make Cowrie logfiles public readable, change the `--umask 0077` option in start.sh into `--umask 0022`

Updating Cowrie

Updating is an easy process. First stop your honeypot. Then fetch updates from GitHub, and upgrade your Python dependencies:

```
bin/cowrie stop
git pull
pip install --upgrade -r requirements.txt
```

If you use output plugins like SQL, Splunk, or ELK, remember to also upgrade your dependencies for these too:

```
pip install --upgrade -r requirements-output.txt
```

And finally, start Cowrie back up after finishing all updates:

```
bin/cowrie start
```

Modifying Cowrie

The pre-login banner can be set by creating the file *honeyfs/etc/issue.net*. The post-login banner can be customized by editing *honeyfs/etc/motd*.

- 2019-10-31 Fix for exec commands when tty logging is disabled
- 2019-10-31 Fix for print output to stdout for curl/wget
- 2019-10-31 Fix for SQL to store full hostname (don't forget to update the database schema)
- 2019-10-15 Slack link now at https://cowrie.org/slack
- 2019-10-04 Subshell ((echo test)) evaluation now working

Release 2.0.0

- 2019-09-06 Crash reporter is enabled by default and will upload data on crashes to api.cowrie.org. This can be disabled in by setting *enabled=false* in *[output_crashreporter]*

- 2019-09-05 Proxy functionality now active by @sgtpepperpt and GSoC2019

- 2019-06-20 Move *auth_none* and *auth_keyboard_interactive_enabled* to [ssh] config section

Release 1.6.0

- 2019-03-31 New documentation theme
- 2019-03-23 Greynoise output plugin (@mzfr)
- 2019-03-19 direct-tcp forwarding now written to databases (@gborges)
- 2019-03-19 Reverse DNS output plugin (@mzfr)
- 2019-03-17 Shell emulation pipe upgrade (@nunonovais)
- 2019-03-14 Shell emulation environment variables improved (@nunonovais)
- 2019-03-14 SSH crypto parameters now configurable in config file (@msharma)
- 2019-03-13 Disable keyboard-interactive authentication by default with option to enable
- 2019-03-13 Added *wc*, *crontab*, *chpasswd* command (@nunonovais)
- 2019-
- 2019-03-07 Output of *ssh -V* now configurable in cowrie.cfg with ssh_version setting
- 2019-03-07 Multiple timezone support in cowrie.cfg timezone directive. Default timezone is now UTC for both cowrie.log and cowrie.json
- 2019-03-12 Handle multiple password prompt. Option to enable or disable keyboard interactive prompt.

Release 1.5.3

- 2019-01-27 Telnet NAWS negotation removed to stop NMAP cowrie detection
- 2019-01-27 Various fixes for Python2/3 compatibility
- 2019-01-09 Documentation converted to ReStructuredText
- 2018-12-04 Fixes for VT outut plugin to only submit new files

CHAPTER 11

Release 1.5.2

- 2018-11-19 Fix tftp exception and tftp test
- 2018-11-14 Remove *dblog* mechanism and *splunk* legacy output plugin.
- 2018-11-01 Add Python3 support for Splunk output plugin
- 2018-10-23 Improved free command
- 2018-10-20 Improved uname command
- 2018-10-16 Save VT results to JSON log

Release 1.5.1

- 2018-10-13 Fixes VT uploads, tab completion on Python3, Hassh support, setuptools functional. userdb migration

- 2018-09-07 NOTE! data/userdb.txt has moved to etc/userdb.txt and a default config is no longer provided!

- 2018-08-25 Downloads and TTY logs have moved to the var/ directory

- 2018-08-11 SSH keys now stored in var/lib/cowrie

- 2018-07-21 source code has move to the src/ directory. Delete old directories twisted/cowrie with compiled code

- 2018-06-29 txtcmds have been moved to share/cowrie/txtcmds

- 2018-06-28 filesystem config entry has changed. please verify if you have custom entry or pickle file

- 2018-06-23 fingerprint log message now holds KEX attributes and a unique fingerprint for the client

- 2018-04-27 Output plugins now require the mandatory config entry 'enabled'.

- 2018-02-06 cowrie.log now uses same rotation mechanism as cowrie.json. One file per day, rather than the default 1MB per file.

- 2017-12-13 Default umask for logs is now 0007. This means group members can access.

- 2017-10-24 Can store uploaded and downloaded artifacts to S3

- 2017-09-23 First proxy implementation for exec commands only

- 2017-07-03 Cuckoo v2 integration

- 2017-05-16 now combines config files: cowrie.cfg.dist and cowrie.cfg in this order

- 2017-05-09 start.sh and stop.sh have been replace by bin/cowrie start|stop

- 2017-04-27 New syntax "listen_endpoints" for configuring listening IP addresses/portnumbers

- 2017-03-15 SSH Forwarding/SFTP/keys/version config have been moved to [ssh]. Change your config file!

- 2017-02-12 Implemented toggle for SSH forwarding

- 2016-08-22 Merged Telnet support by @obilodeau!

- 2016-08-20 Update your libraries! 'configparser' now required: "pip install configparser"

- 2016-05-06 Load pickle once at startup for improved speed

- 2016-04-28 files in utils/ have been moved to bin/

- 2016-01-19 Support openssh style delayed compression

- 2016-01-13 Correct '.' support and +s and +t bits in ls

- 2016-01-13 Full username/group in SFTP ls

- 2016-01-05 Basic VirusTotal support has been added

- 2016-01-04 No longer crash when client tries ecdsa

- 2015-12-28 Interact port (default 5123) only listens on loopback interface now (127.0.0.1)

- 2015-12-24 Redirect to file (>) now works for most commands and is logged in dl/ directory

- **2015-12-06 UID information is now retrieved from honeyfs/etc/passwd. If you added additional users** you will need to add these to the passwd file as well

- 2015-12-04 New 'free' command with '-h' and '-m' options

- 2015-12-03 New 'env' command that prints environment variables

- 2015-02-02 Now use honeyfs/etc/passwd and group to get uid/gid info

- 2015-11-29 Size limit now enforced for SFTP uploads

- 2015-11-25 New 'sudo' command added

- **2015-11-19 Queued input during commands is now sent to shell to be executed** when command is finished

- 2015-11-18 Added SANS DShield output (Thanks @UnrealAkama)

- 2015-11-17 Added ElasticSearch output (Thanks @UnrealAkama)

- 2015-11-17 Standard input is now saved with SHA256 checksum. Duplicate data is not saved

- 2015-11-12 New 'busybox' command added (Thanks @mak)

- **2015-09-26 keyboard-interactive is back as authentication method, after** Twisted removed support initially

- 2015-07-30 Local syslog output module

- 2015-06-15 Cowrie now has a '-c' startup switch to specify the configuration file

- 2015-06-15 Removed exec_enabled option. This feature is now always enabled

- 2015-06-03 Cowrie now uses twisted plugins and has gained the '-p' commandline option

- 2015-06-01 Cowrie no longer search for config files in /etc and /etc/cowrie

- 2015-04-12 JSON output is now default via 'output' plugin mechanism. Rotates daily

- 2015-04-10 Fix for downloading files via SFTP

- 2015-03-31 Small tweaks on session close, closing session does not close ssh transport

- 2015-03-18 Merged 'AuthRandom' login class by Honigbij

- **2015-02-25 Internals for dblog/ modules changed completely.** Now accepts structured logging arguments, and uses eventids instead of regex parsing

- 2015-02-20 Removed screen clear/reset on logout

- 2015-02-19 Configuration directives have changed! ssh_addr has become listen_addr and ssh_port has become listen_port. The old keywords are still accepted for backwards compatibility

- default behaviour is changed to disable the exit jail

- sftp support

- exec support

- **stdin is saved as a file in dl/ when using exec commands** to support commands like 'cat >file; ./file'

- allow wget download over non-80 port

- simple JSON logging added

- accept log and deny publickey authentication

- add uname -r, -m flags

- add working sleep command

- enabled ssh diffie-hellman-group-exchange-sha1 algorithm

- add 'bash -c' support (no effect option)

- enable support for && multiple commands

- create uuid to uniquely identify each session

- log and deny direct-tcpip attempts

- add "chattr" command

- support emacs keybindings (c-a, c-b, c-f, c-p, c-n, c-e)

- add "sync" command

- accept, log and deny public key authentication

- add "uname -r" support

- logstash and kibana config files added, based on JSON log

- fix for honeypot detection (pre-auth differences with openssh)

- added verbose logging of client requested key exchange parameters (for client fingerprinting)

- fixes for behavior with non-existent files (cd /test, cat /test/nonexistent, etc)

- fix for ability to ping/ssh non-existent IP address

- always send ssh exit-status 0 on exec and shell

- ls output is now alphabetically sorted

- banner_file is deprecated. honeyfs/etc/issue.net is default

- add 'dir' alias for 'ls'

- add 'help' bash builtin

- add 'users' aliased to 'whoami'

- add 'killall' and 'killall5' aliased to nop

- add 'poweroff' 'halt' and 'reboot' aliases for shutdown

- add environment passing to commands

- added 'which', 'netstat' and 'gcc' from kippo-extra

- logging framework allows for keyword use

Contributing Guidelines

Thank you for your interest in contributing to our project. Whether it's a bug report, new feature, correction, or additional documentation, we greatly value feedback and contributions from our community.

Please read through this document before submitting any issues or pull requests to ensure we have all the necessary information to effectively respond to your bug report or contribution.

Reporting Bugs/Feature Requests

We welcome you to use the GitHub issue tracker to report bugs or suggest features.

When filing an issue, please check existing open, or recently closed, issues to make sure somebody else hasn't already reported the issue. Please try to include as much information as you can. Details like these are incredibly useful:

- A reproducible test case or series of steps

- The version of our code being used

- Any modifications you've made relevant to the bug

- Anything unusual about your environment or deployment

Contributing via Pull Requests

Contributions via pull requests are much appreciated. Before sending us a pull request, please ensure that:

1. You are working against the latest source on the *master* branch.

2. You check existing open, and recently merged, pull requests to make sure someone else hasn't addressed the problem already.

3. You open an issue to discuss any significant work - we would hate for your time to be wasted.

To send us a pull request, please:

1. Fork the repository.

2. Modify the source; please focus on the specific change you are contributing. If you also reformat all the code, it will be hard for us to focus on your change.

3. Ensure local tests pass.

4. Commit to your fork using clear commit messages.

5. Send us a pull request, answering any default questions in the pull request interface.

6. Pay attention to any automated CI failures reported in the pull request, and stay involved in the conversation.

GitHub provides additional document on forking a repository and creating a pull request.

Finding contributions to work on

Looking at the existing issues is a great way to find something to contribute on. As our projects, by default, use the default GitHub issue labels ((enhancement/bug/duplicate/help wanted/invalid/question/wontfix), looking at any help wanted' issues is a great place to start.

Licensing

See the LICENSE file for our project's licensing. We will ask you confirm the licensing of your contribution.

Using the Proxy

The SSH and Telnet proxies can be used to provide a fully-fledged environment, in contrast to the emulated shell traditionally provided by Cowrie. With a real backend environment where attackers can execute any Unix command, Cowrie becomes a high-interaction honeypot.

To use the proxy, start by changing the *backend* option to *proxy* in the *[honeypot]* section. In the remainder of this guide we will refer to the *[proxy]* section of the config file.

18.1 Choosing a Backend

Cowrie supports a simple backend (i.e., a real machine or virtual machines provided by you), but you can use Cowrie's backend pool, which provides a set of VMs, handling their boot and cleanup, also ensuring that different attackers (different IPs) each see a "fresh" environment, while connections from the same IP get the same VM.

VERY IMPORTANT NOTE: some attacks consist of downloading malicious software or accessing illegal content through insecure machines (such as your honeypot). If you are using your **own backend**, be sure to restrict networking to the Internet on your backend, and ensure other machines on your local network are isolated from the backend machine. The backend pool restricts networking and does its best to ensure total isolation, to the best of Qemu/libvirt (and our own) capabilities. **Be very careful to protect your network and devices!**

18.2 Configuring the Proxy

18.2.1 Backend configs

If you choose the simple backend, configure the hosts and ports for your backend. For the backend pool, configure the variables starting with *pool_*. You'll also need to deal with the *[backend_pool]* section, which we detail in the Backend Pool's own documentation.

The backend pool can be run in the same machine as Cowrie, or on a remote one (e.g. Cowrie on a Raspberry Pi, and the pool in a larger machine). In the former case, set *pool* to *local*; in the later, set *pool* to *remote* and specify its host

and port, matching with the *listen_endpoints* of the *[backend_pool]* section. Further configurations sent by the client are explained in Backend Pool's own documentation.

18.2.2 Authentication

Regardless of the used type of backend, Cowrie will need credentials to access the machine. These can be of any account on it, as long as it supports password authentication.

Note that these are totally independent of the credentials attackers can use (as set in *userdb*). *userdb* credentials are the ones attackers may use to connect to Cowrie, while *backend_user* and *backend_pass* are used to connect Cowrie to the backend.

18.2.3 Telnet prompt detection

Due to the different implementations of Telnet, there is not a single reliable way of catching the authentication phase of the protocol as in SSH. Therefore, we rely on regex expressions to detect authentication prompts, allowing us to identify the credentials supplied by the attacker and check if they are accepted by *userdb*. If they are, we send the *backend_user* and *backend_pass* to the backend (spoofing the authentication); if not, we send *backend_pass* appended with the word *fake* to force a login failed prompt (and fail authentication overall).

If you don't want to spoof authentication, set *telnet_spoof_authentication* to false. In this mode, only the backend real details will be accepted to authenticate, thus bypassing *userdb*.

The expressions to detect authentication prompts are *telnet_username_prompt_regex* and *telnet_password_prompt_regex*. A further expression we use is defined in *telnet_username_in_negotiation_regex*. Some clients send their username in the first phases of the protocol negotiation, which some systems (the backend) use to only show the password prompt the first time authentication is tried (thus assuming the client's username as the username they'll use to login into the system). Cowrie tries to capture this username and use it when comparing the auth details with the *userdb*.

18.2.4 Analysing traffic

Analysing raw traffic can be interesting when setting up Cowrie, in particular to set-up Telnet prompt detection. For this, you can set *log_raw* to true.

Changing the Cowrie file system

19.1 Introduction

Part of Cowrie is an emulated file system. Each honeypot visitor will get their own personal copy of this file system and this will deleted when they log off. They can delete or change any file, nothing will be preserved.

The file system implementation consists of two pieces, the *pickle* file, which holds metadata for the files (filename, directory, permissions, owner, size, file type, etc), and the *honeyfs* directory that holds file contents. The honeyfs directory only has certain files by default. Most files do not have content associated with them.

To show the contents of the file, it needs both a meta data entry (pickle) and a honeyfs file.

19.2 Creating a new pickle file

Create a directory where you put all files you'd like to be show in your filesystem Create the pickle file:

```
./bin/createfs -l YOUR-DIR -d DEPTH -o ./share/cowrie/custom.pickle
```

Make sure your config picks up custom.pickle, or rename it to fs.pickle

The pickle file just gives the layout of filesystem. If you actually need the files to be really there (for example to read them), you'll have to create them in honeyfs/xxx

Backend Pool

The Backend Pool manages a set of dynamic backend virtual machines to be used by Cowrie's proxy. The pool keeps a set of VMs running at all times, ensuring different attackers each see a "pristine" VM, while repeated connections from the same IP are served with the same VM, thus ensuring a consistent view to the attacker. Furthermore, VMs in the pool have their networking capabilities restricted by default: some attacks consist of downloading malicious software or accessing illegal content through insecure machines (such as your honeypot). Therefore, we limit any access to the Internet via a network filter, which you can configure as you see fit.

The VMs in the backend pool, and all infrastructure (snapshots, networking and filtering) are backed-up by Qemu/libvirt. We provide two example VM images (for Ubuntu Server 18.04 and OpenWRT 18.06.4) whose configurations are already set and ready to be deployed. Further below in this guide we'll discuss how to create your own images and customise libvirt's XML configuration files.

First of all, install the needed dependencies for the pool, as explained in the installation steps.

20.1 Provided images

To allow for a simple setup, we provide two VM images to use with the backend pool: Ubuntu 18.04 and OpenWRT. You can download them below, and then edit *cowrie.cfg* to match the path of the images. In the case of OpenWRT you will need two different files. Note that a separate set of configs is provided for each image in the default configuration. Choose the one you want to use and comment the other as needed.

- Ubuntu 18.04.
- OpenWRT disk image.
- OpenWRT kernel image.

20.2 Backend Pool initialisation

Depending on the machine that will be running the backend pool, initialisation times for VMs can vary greatly. If the pool is correctly configured, you will get the *PoolServerFactory starting on 6415* message on your log.

After a while, VMs will start to boot and, when ready to be used, a message of the form *Guest 0 ready for connections @ 192.168.150.68! (boot 17s)* will appear for each VM. Before VMs are ready, SSH and Telnet connections from attackers will be dropped by Cowrie.

20.3 Proxy configurations

When the proxy starts, and regardless whether the backend pool runs on the same machine as the proxy or not, some configurations are sent by the proxy to the pool during runtime.

These are:

- **pool_max_vms**: the number of VMs to be kept running in the pool

- **pool_vm_unused_timeout**: how much time (seconds) a used VM is kept running (so that an attacker that reconnects is served the same VM.

- **apool_share_guests**: what to do if no "pristine" VMs are available (i.e., all have been connected to); if set to true we serve a random one from the used, if false we throw an exception.

20.4 Backend Pool configuration

In this section we'll discuss the *[backend_pool]* section of the configuration file.

The backend pool can be run in the same machine as the rest of Cowrie, or in a separate one. In the former case, you'd be running Cowrie with

```
[backend_pool]
pool_only = false

[proxy]
backend = pool
pool = local
```

If you want to deploy the backend pool in a different machine, then you'll need to invert the configuration: the pool machine has *pool_only = true* (SSH and Telnet are disabled), and the proxy machine has *pool = remote*.

Note: The communication protocol used between the proxy and the backend pool is unencrypted. Although no sensitive data should be passed, we recommend you to only use private or local interfaces for listening when setting up *listen_endpoints*.

20.4.1 Recycling VMs

Currently, handling of virtual machines by the pool is not perfect. Sometimes, VMs reach an inconsistent state or become unreliable. To counter that, and ensure fresh VMs are ready constantly, we use the *recycle_period* to periodically terminate running instances, and boot new ones.

20.4.2 Snapshots

VMs running in the pool are based on a base image that is kept unchanged. When booting, each VM creates a snaphost that keeps track of differences between the base image and snapshot. If you want to analyse snapshots and see any changes made in the VMs, set *save_snapshots* to true. If set to true be mindful of space concerns, each new VM will take at least ~20MB in storage.

20.4.3 XML configs (advanced)

You can change libvirt's XML configs from the default ones in *share/cowrie/pool_configs*. However, if you're using one of the default images provided, then you probably don't need to.

20.4.4 Guest configurations

A set of guest (VM) parameters can be defined as we explain below:

- **guest_config**: the XML configuration for the guest (default_guest.xml works for x86 machines, and wrt_arm_guest.xml for ARM-based OpenWRT)

- **guest_privkey**: currently unused

- **guest_tag**: an identifiable name for snapshots and logging

- **guest_ssh_port / guest_telnet_port**: which ports are listening for these on the VM (no relation with the ports Cowrie's listening to)

- **guest_image_path**: the base image upon which all VMs are created from

- **guest_hypervisor**: the hypervisor used; if you have an older machine or the emulated architecture is different from the host one, then use software-based "qemu"; however, if you are able to, use "kvm", it's **much** faster.

- **guest_memory**: memory assigned to the guest; choose a value considering the number of guests you'll have running in total (*pool_max_vms*)

20.4.5 NATing

VMs are assigned an IP in a local network defined by libvirt. If you need to access the VMs from a different machine (i.e., running the backend pool remotely), then an external-facing IP (as defined in *nat_public_ip*) is needed for the proxy to connect to.

For this purpose, we provide a simple form of NAT that, for each VM request, and if enabled, starts a TCP proxy to forward data from a publicly-acessible IP to the internal libvirt interface.

20.5 Creating VM images

Creating a new type of VM involves three steps: creating a base image, installing the OS, and tweaking configs.

To create a disk image issue

```
sudo qemu-img create -f qcow2 image-name.qcow2 8G
```

(the qcow2 format is needed to ensure create snapshots, thus providing isolation between each VM instance; you can specify the size you want for the disk)

Then you'll have to install an OS into it

```
virt-install --name temp-domain --memory 1024 --disk image-name.qcow2 --cdrom os-
→install-cd.iso --boot cdrom
```

(to use virt-install you need to install the virtinst package)

After install check that the VM has network connectivity. If you set the pool to use the created image and SSH does not come up, log into the VM via libvirt (e.g., using virt-manager) and try the following (might change depending on system)

```
# run all as root
ip link show
ip link set enp1s0 up
dhclient
```

In Ubuntu dhclient can be set to run with netplan, for example, to be run on startup.

20.5.1 Set up Telnet

Steps used in Ubuntu, can be useful in other distros

```
# run all as root
apt-get -y install telnetd xinetd
touch /etc/xinetd.d/telnet
printf "service telnet\n{\ndisable = no\nflags = REUSE\nsocket_type = stream\nwait =␣
→no\nuser = root\nserver = /usr/sbin/in.telnetd\nlog_on_failure += USERID\n}" > /etc/
→xinetd.d/telnet
printf "pts/0\npts/1\npts/2\npts/3\npts/4\npts/5\npts/6\npts/7\npts/8\npts/9" >> /etc/
→securetty
service xinetd start
```

20.5.2 Customising XML configs

If you want, you can customise libvirt's XML configurations.

The main configuration for a guest is located in *default_guest.xml*. This defines the virtual CPU, available memory, and devices available on the guest. Most of these configurations are set by Cowrie using the guest configurations; you'll see them in the XML as templates ("{guest_name}"). The main blocks of XML regard the disk and network interface devices.

You can include these blocks as-is in your custom configuration to allow Cowrie to manage your custom guests automatically.

```
<disk type='file' device='disk'>
    <driver name='qemu' type='qcow2'/>
    <source file='{disk_image}'/>
    <target dev='vda' bus='virtio'/>
    <address type='pci' domain='0x0000' bus='0x03' slot='0x00' function='0x0'/>
</disk>
```

```
<interface type='network'>
    <start mode='onboot'/>
    <mac address='{mac_address}'/>
    <source network='{network_name}'/>
    <model type='virtio'/>
    <address type='pci' domain='0x0000' bus='0x01' slot='0x00' function='0x0'/>
    <filterref filter='cowrie-default-filter'/>
</interface>
```

The other important configuration file is *default_filter.xml*, which handles how networking is restricted in the guest VM (aka to the attackers). This file is composed by a set of rules of the form

```
<rule action='accept' direction='in'>
    <tcp dstportstart='22'/>
</rule>
```

Each rule specifies a type of traffic (TCP, UDP. . .) and direction, whether to accept or drop that traffic, and the destiantion of traffic. The default filter provided allows inbound SSH and Telnet connections (without which the VM would be unusable, outbound ICMP traffic (to allow pinging) and outbound DNS querying. All other traffic is dropped as per the last rule, thus forbidding any download or tunnelling.

VERY IMPORTANT NOTE: some attacks consist of downloading malicious software or accessing illegal content through insecure machines (such as your honeypot). Our provided filter restricts networking and does its best to ensure total isolation, to the best of Qemu/libvirt (and our own) capabilities. **Be very careful to protect your network and devices while allowing any more traffic!**

20.6 References

- libvirt guest XML syntax
- libvirt network filter XML syntax
- Create a OpenWRT image

Analysing snapshots and downloaded content

One interesting aspect of Cowrie is the capability to analyse any downloaded malware and content into the honeypot. The snapshot mechanism can be leveraged to analyse any download and any change performed against the base image, to determine which files have been changed.

This guide shows how that can be achieved by leveraging using the *libguestfs-tools* package.

21.1 Getting the a filesystem diff

The first step is getting the differences between each VM that was used and the base image provided. The tool we'll be using is *virt-diff*, which provides a similar syntax to that of Unix's *diff*.

```
sudo virt-diff -a ~/cowrie-imgs/ubuntu18.04-minimal.qcow2 -A snapshot-ubuntu18.04-
↪a70b9671ad4d44619af2c4a41a28aec0.qcow2
```

(the tool might need to be run with sudo due to a permission denied error)

The output will contain all changed files and their content, which might get long easily. The following command outputs the names of changed files, to be easier to read (assuming the output from *virt-diff* is stored in a file diff.txt)

```
grep -aE "^\+ |^- |^= " diff.txt
```

Here is an example output, in a VM were we created a file called *avirus*:

```
= - 0644      1024 /boot/grub/grubenv
= - 0600      1036 /root/.bash_history
+ - 0644        14 /root/avirus
+ d 0700      4096 /tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪resolved.service-syUmHS
+ d 1777      4096 /tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪resolved.service-syUmHS/tmp
+ d 0700      4096 /tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪timesyncd.service-SrDysr
```

(continues on next page)

```
+ d 1777       4096 /tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪timesyncd.service-SrDysr/tmp
+ - 0644       2163 /var/backups/apt.extended_states.0
+ - 0644          0 /var/lib/apt/daily_lock
= - 0644          0 /var/lib/private/systemd/timesync/clock
= - 0600        512 /var/lib/systemd/random-seed
= - 0644          0 /var/lib/systemd/timers/stamp-apt-daily-upgrade.timer
= - 0644          0 /var/lib/systemd/timers/stamp-apt-daily.timer
= - 0644          0 /var/lib/systemd/timers/stamp-fstrim.timer
= - 0644          0 /var/lib/ubuntu-release-upgrader/release-upgrade-available
= - 0640      14534 /var/log/auth.log
= - 0640    8388608 /var/log/journal/19497399992e49388d57aa395b993b2c/system.journal
= - 0640     346383 /var/log/kern.log
= - 0664     292292 /var/log/lastlog
= - 0640     436458 /var/log/syslog
= - 0664      19968 /var/log/wtmp
+ d 0700       4096 /var/tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪resolved.service-u5dZk6
+ d 1777       4096 /var/tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪resolved.service-u5dZk6/tmp
+ d 0700       4096 /var/tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪timesyncd.service-Tcil4E
+ d 1777       4096 /var/tmp/systemd-private-9f5f6c41f75f48f4991c55f3fc3d6435-systemd-
↪timesyncd.service-Tcil4E/tmp
```

As you can see, the created file is shown among lots of log and temporary files. There is no good way to eliminate these, but we can use grep to ignore them:

```
grep -aE "^\+ |^- |^= " diff.txt | grep -aEv "/tmp/systemd|/var/log|/var/lib"
```

Which now gives us a clearer output:

```
= - 0644       1024 /boot/grub/grubenv
= - 0600       1036 /root/.bash_history
+ - 0644         14 /root/avirus
+ - 0644       2163 /var/backups/apt.extended_states.0
```

21.2 Getting interesting files

To be able to get and read the files you're interested in, you'll need to mount the snapshot into your machine and copy the file(s) into your disk. The steps we describe are taken from here, and rewritten here for clarity.

We start by mounting the image in a temporary dir:

```
mkdir /tmp/mount_qcow2
sudo guestmount -a snapshot-ubuntu18.04-a70b9671ad4d44619af2c4a41a28aec0.qcow2 -m /
↪dev/sda1 --ro /tmp/mount_qcow2
```

If we now search for the file in the mount directory we can see its contents, and then unmount the drive:

```
sudo ls -halt /tmp/mount_qcow2/root
total 32K
-rw-------  1 root root 1.1K Jul 28 21:45 .bash_history
drwx------  3 root root 4.0K Jul 28 21:45 .
```

```
-rw-r--r--  1 root root   14 Jul 28 21:45 avirus
drwxr-xr-x 22 root root 4.0K Jul 15 01:57 ..
-rw-r--r--  1 root root   74 Jul 15 00:59 .selected_editor
drwx------  2 root root 4.0K Jul 15 00:59 .cache
-rw-r--r--  1 root root 3.1K Apr  9  2018 .bashrc
-rw-r--r--  1 root root  148 Aug 17  2015 .profile

sudo cat /tmp/mount_qcow2/root/avirus
virus content

sudo guestunmount /tmp/mount_qcow2/
```

Note: the device to be mounted from the image isn't always */dev/sda1*. However, if you run the command as-is, *guestmount* will check if */dev/sda1* exists and, if not, it will list available partitions for you.

How to process Cowrie output in an ELK stack

(Note: work in progress, instructions are not verified)

22.1 Prerequisites

- Working Cowrie installation
- Cowrie JSON log file (enable database json in cowrie.cfg)
- Java 8

22.2 Installation

This is a simple setup for ELK stack, to be done on the same machine that is used for cowrie. We use *Filebeat* to send logs to *Logstash*, and we use *Nginx* as a reverse proxy to access *Kibana*. Note there are many other possible configurations!

Add Elastic's repository and key:

```
wget -qO - https://packages.elastic.co/GPG-KEY-elasticsearch | sudo apt-key add -
echo "deb https://artifacts.elastic.co/packages/7.x/apt stable main" | sudo tee /etc/
→apt/sources.list.d/elastic-7.x.list
apt-get update
```

Install logstash, elasticsearch, kibana and filebeat:

```
sudo apt -y install apt-transport-https wget default-jre
sudo apt install elasticsearch logstash kibana
sudo apt install filebeat
sudo apt install nginx apache2-utils
```

Enable the services:

```
sudo systemctl enable elasticsearch logstash kibana filebeat nginx
```

22.3 ElasticSearch Configuration

ElasticSearch configuration file is located in */etc/elasticsearch/elasticsearch.yml*. The default settings need not be changed.

By default, ElasticSearch listens on port 9200. Test it:

```
curl http://localhost:9200
```

You should get a JSON object in return.

22.4 Kibana Configuration

Make a folder for logs:

```
sudo mkdir /var/log/kibana
sudo chown kibana:kibana /var/log/kibana
```

Change the following parameters in */etc/kibana/kibana.yml* to reflect your server setup:

```
"server.host"  - set it to "localhost" if you use nginx for basic authentication or
→external interface if you use XPack (see below)
"server.name" - name of the server
"elasticsearch.url" - address of the elasticsearch: ["http://localhost:9200"]
"elasticsearch.username", "elasticsearch.password" - needed only if you use XPack
→(see below)
"logging.dest" - set path to logs (/var/log/kibana/kibana.log)
```

22.5 Logstash Configuration

Get GeoIP data from www.maxmind.com (free but requires registration): download the GeoLite2 City GZIP. Unzip it and locate the mmdb file. Place it somewhere in your filesystem and make sure that "logstash" user can read it:

```
sudo mkdir -p /opt/logstash/vendor/geoip/
sudo mv GeoLite2-City.mmdb /opt/logstash/vendor/geoip
```

Configure logstash:

```
sudo cp logstash-cowrie.conf /etc/logstash/conf.d
```

Make sure the configuration file is correct. Check the input section (path), filter (geoip databases) and output (elastic-search hostname):

```
sudo systemctl restart logstash
```

22.6 Filebeat Configuration

Filebeat is not mandatory (it is possible to directly read cowrie logs from Logstash) but nice to have, because if Logstash is under pressure, it automatically knows to slow down + it is possible to deal with multiple sensor inputs.

Configure filebeat:

```
sudo cp filebeat-cowrie.conf /etc/filebeat/filebeat.yml
```

Check the following parameters:

```
filebeat.inputs: the path must point to cowrie's json logs
output.elasticsearch: must be false because we want Filebeat to send to Logstash, not
→directly to ElasticSearch
output.logstash: must be true. The default port for Logstash is 5044, so hosts should
→be ["localhost:5044"]
```

Start filebeat:

```
sudo systemctl start filebeat
```

22.7 Nginx

ELK has been configured on localhost. If you wish to access it remotely, you can setup a reverse proxy to Kibana's backend server, which runs on port 5601 by default.

Install Nginx:

```
sudo apt install nginx apache2-utils
```

Create an administrative Kibana user and password:

```
sudo htpasswd -c /etc/nginx/htpasswd.users admin_kibana
```

Edit Nginx configuration /etc/nginx/sites-available/default. Customize port to what you like, and specify your server's name (or IP address):

```
server {
    listen YOURPORT;

    server_name YOURIPADDRESS;

    auth_basic "Restricted Access";
    auth_basic_user_file /etc/nginx/htpasswd.users;

    location / {
        proxy_pass http://localhost:5601;
        proxy_http_version 1.1;
        proxy_set_header Upgrade $http_upgrade;
        proxy_set_header Connection 'upgrade';
        proxy_set_header Host $host;
        proxy_cache_bypass $http_upgrade;
    }
}
```

Start the service:

```
sudo systemctl start nginx
```

22.8 Using Kibana

You can list indexes with:

```
curl 'http://localhost:9200/_cat/indices?v'
```

You should see a cowrie index cowrie-logstash-DATE. . . Its health is yellow because the number of replicas should be set to 0 (unless you want another configuration):

```
curl -XPUT 'localhost:9200/cowrie-logstash-REPLACEHERE/_settings' -H "Content-Type:␣
→application/json" -d '{ "index" : {"number_of_replicas" : 0 } }'
```

It should answer {"acknowledged":true}

In Kibana's GUI, create an index pattern (Management / Index Patterns) for

```
cowrie-logstash-*
```

Use default settings and timestamp.

22.9 Tuning ELK stack

Refer to elastic's documentation about proper configuration of the system for the best elasticsearch's performance

You may avoid installing nginx for restricting access to kibana by installing official elastic's plugin called "X-Pack" (https://www.elastic.co/products/stack)

ELK log files get big: ensure you have enough space in /var, consider setting up LVM or ZFS partitions.

22.10 Troubleshooting

- View service logs with: sudo journalctl -u service
- If the date in Kibana is incorrect, check (Advanced Settings / dateFormat)

How to process Cowrie output into Graylog

23.1 Prerequisites

- Working Cowrie installation
- Working Graylog installation

23.2 Cowrie Configuration

Open the Cowrie configuration file and uncomment these 3 lines:

```
[output_localsyslog]
facility = USER
format = text
```

Restart Cowrie

23.3 Graylog Configuration

Open the Graylog web interface and click on the **System** drop-down in the top menu. From the drop-down menu select **Inputs**. Select **Syslog UDP** from the drop-down menu and click the **Launch new input** button. In the modal dialog enter the following information:

```
**Title:** Cowrie
**Port:** 8514
**Bind address:** 127.0.0.1
```

Then click **Launch.**

23.4 Syslog Configuration

Create a rsyslog configuration file in /etc/rsyslog.d:

```
$ sudo nano /etc/rsyslog.d/85-graylog.conf
```

Add the following lines to the file:

```
$template GRAYLOGRFC5424,"<%pri%>%protocol-version% %timestamp:::date-rfc3339%
→%HOSTNAME% %app-name% %procid% %msg%\n"
*.* @127.0.0.1:8514;GRAYLOGRFC5424
```

Save and quit.

Restart rsyslog:

```
$ sudo service rsyslog restart
```

How to process Cowrie output in kippo-graph

(Note: work in progress, instructions are not verified) Tested on Debian 9.

24.1 Prerequisites

- Working Cowrie installation
- LAMP stack (Linux, Apache, MySQL, PHP)

24.2 Installation

This covers a simple installation, with kippo-graph and Cowrie on the same server. Please see here for installation: https://github.com/ikoniaris/kippo-graph

24.3 MySQL configuration

Configuring Cowrie requires setting up the SQL tables and then telling Cowrie to use them.

To install the tables and create the Cowrie user account enter the following commands:

```
mysql -u root -p
CREATE DATABASE cowrie;
GRANT ALL ON cowrie.* TO 'cowrie'@'localhost' IDENTIFIED BY 'PASSWORD HERE';
FLUSH PRIVILEGES;
exit
```

Next create the database schema:

```
cd /opt/cowrie/
mysql -u cowrie -p
USE cowrie;
source ./docs/sql/mysql.sql;
exit
```

disable MySQL strict mode:

```
vi /etc/mysql/conf.d/disable_strict_mode.cnf

[mysqld]
sql_mode=IGNORE_SPACE,NO_ZERO_IN_DATE,NO_ZERO_DATE,ERROR_FOR_DIVISION_BY_ZERO,NO_AUTO_
→CREATE_USER,NO_ENGINE_SUBSTITUTION
```

24.4 Cowrie configuration

Edit cowrie.cfg:

```
vi /opt/cowrie/cowrie.cfg
```

Activate output to mysql:

```
[output_mysql]
host = localhost
database = cowrie
username = cowrie
password = PASSWORD HERE
port = 3306
debug = false
```

Set read access to tty-files for group www-data (group maybe differ on other distributions):

```
sudo apt-get install acl
sudo setfacl -Rm g:www-data:rx /opt/cowrie/var/lib/cowrie/tty/
```

24.5 kippo-graph Configuration

Edit config file:

```
vi /var/www/html/kippo-graph/config.php
```

Change db settings:

```
define('DB_HOST', 'localhost');
define('DB_USER', 'cowrie');
define('DB_PASS', 'PASSWORD HERE');
define('DB_NAME', 'cowrie');
define('DB_PORT', '3306');
```

24.6 Apache2 configuration (optional)

To secure the installation

Create password database:

```
cd /etc/apache2/
htpasswd -c /etc/apache2/cowrie.passwd <username>
htpasswd /etc/apache2/cowrie.passwd <username> (second user)


vi /etc/apache2/sites-enabled/000-default.conf
```

Between the <VirtualHost> </VirtualHost> tags, add:

```
<Location />
    AuthBasicAuthoritative On
    AllowOverride AuthConfig

    AuthType Basic
    AuthName "cowrie honeypot"
    AuthUserFile /etc/apache2/cowrie.passwd
    Require valid-user
</Location>
```

How to process Cowrie output with Splunk

25.1 Splunk Output Module

- In Splunk, enable the HTTP Event Collector (go to Settings->Add Data)
- Do not enable *Indexer Acknowledgment*
- Copy the authorization token for later use
- Modify *cowrie.cfg* to enable the *[splunk]* section
- Add URL to HTTP Event Collector and add the authorization token
- Optionally enable sourcetype, source, host and index settings

25.2 File Based

- Collect cowrie.json output file using Splunk

25.3 Reporting

Please see: https://github.com/aplura/Tango

How to Send Cowrie Output to a MySQL Database

26.1 Prerequisites

- Working Cowrie installation
- MySQL Server installation

26.2 Installation

Run:

```
$ sudo apt-get install mysql-server libmysqlclient-dev python-mysqldb
$ su - cowrie
$ source cowrie/cowrie-env/bin/activate
$ pip install mysqlclient
```

Previously MySQL-python was used. Only if you run into isses with mysqlclient, try this instead:

```
$ pip install MySQL-python
```

26.3 MySQL Configuration

First create an empty database named 'cowrie':

```
$ mysql -u root -p
CREATE DATABASE cowrie;
```

Create a cowrie user account for the database and grant all access privileges:

```
GRANT ALL ON cowrie.* TO 'cowrie'@'localhost' IDENTIFIED BY 'PASSWORD HERE';
```

Restricted Privileges:

Alternatively you can grant the cowrie account with less privileges. The following command grants the account with the bare minimum required for the output logging to function:

```
GRANT INSERT, SELECT, UPDATE ON cowrie.* TO 'cowrie'@'localhost' IDENTIFIED BY
↪'PASSWORD HERE';
```

Apply the privilege settings and exit mysql:

```
FLUSH PRIVILEGES;
exit
```

Next, log into the MySQL database using the cowrie account to verify proper access privileges and load the database schema provided in the docs/sql/ directory:

```
$ cd ~/cowrie/docs/sql/
$ mysql -u cowrie -p
USE cowrie;
source mysql.sql;
exit
```

26.4 Cowrie Configuration

Uncomment and update the following entries to ~/cowrie/cowrie.cfg under the Output Plugins section:

```
[output_mysql]
host = localhost
database = cowrie
username = cowrie
password = PASSWORD HERE
port = 3306
debug = false
enabled = true
```

Restart Cowrie:

```
$ cd ~/cowrie/bin/
$ ./cowrie restart
```

Verify That the MySQL Output Engine Has Been Loaded

Check the end of the ~/cowrie/log/cowrie.log to make sure that the MySQL output engine has loaded successfully:

```
$ cd ~/cowrie/log/
$ tail cowrie.log
```

Example expected output:

```
2017-11-27T22:19:44-0600 [-] Loaded output engine: jsonlog
2017-11-27T22:19:44-0600 [-] Loaded output engine: mysql
...
2017-11-27T22:19:58-0600 [-] Ready to accept SSH connections
```

Confirm That Events are Logged to the MySQL Database Wait for a new login attempt to occur. Use tail like before to quickly check if any activity has been recorded in the cowrie.log file.

Once a login event has occurred, log back into the MySQL database and verify that the event was recorded:

```
$ mysql -u cowrie -p
USE cowrie;
SELECT * FROM auth;
``
```

Example output:

```
+----+-------------+---------+---------+------------+---------------------+
| id | session     | success | username | password  | timestamp           |
+----+-------------+---------+---------+------------+---------------------+
|  1 | a551c0a74e06 |       0 | root    | 12345      | 2017-11-27 23:15:56 |
|  2 | a551c0a74e06 |       0 | root    | seiko2005  | 2017-11-27 23:15:58 |
|  3 | a551c0a74e06 |       0 | root    | anko       | 2017-11-27 23:15:59 |
|  4 | a551c0a74e06 |       0 | root    | 123456     | 2017-11-27 23:16:00 |
|  5 | a551c0a74e06 |       0 | root    | dreambox   | 2017-11-27 23:16:01 |
...
```

Using TCP tunneling with Squid

27.1 Prerequisites

- Working Cowrie installation
- Working Squid installation with CONNECT allowed
- (optional) Rate limit and black/white lists in Squid

27.2 Installation

`` $ sudo apt-get install squid ``

27.3 Squid Configuration

See *squid.conf* for an example configuration.

27.4 Cowrie Configuration

Uncomment and update the following entries to ~/cowrie/cowrie.cfg under the SSH section:

`` forward_tunnel = true

forward_tunnel_80 = 127.0.0.1:3128 forward_tunnel_443 = 127.0.0.1:3128 ``

Restart Cowrie

`` $ cd ~/cowrie/bin/ $ ./cowrie restart ``

Automatically starting Cowrie with supervisord

- Copy the file *cowrie.conf* to */etc/supervisor/conf/*

Automatically starting Cowrie with systemd

NOTE: untested

- Copy the file *docs/systemd/system/cowrie.socket* to */etc/systemd/system*

- Copy the file *docs/systemd/system/cowrie.service* to */etc/systemd/system*

- Examine */etc/systemd/system/cowrie.service* and ensure the paths are correct for your installation if you use non-standard file system locations.

- Add entries to *etc/cowrie.cfg* to listen on ports via systemd. These must match your cowrie.socket configuration:

    [ssh] listen_endpoints = systemd:domain=INET6:index=0

    [telnet] listen_endpoints = systemd:domain=INET6:index=1

- **Run:** sudo systemctl daemon-reload sudo systemctl start cowrie.service sudo systemctl enable cowrie.service

CHAPTER 30

Indices and tables

- genindex
- modindex
- search